

AI Metadata and Session Data Repository: A Client Architecture for Managing, Measuring, and Mitigating Enterprise Generative AI

Johnny Morgan
UMBC Department of Information Systems
yj07538@umbc.edu

August 13, 2026

Abstract

Enterprises are deploying generative AI into business processes without the instrumentation needed to manage it. We propose a client architecture that captures AI utilization metadata and session data as the precondition for managing AI systems: an inventory of AI tools and instances, quantitative and qualitative measures of utilization and effectiveness (KPIs, MOEs, MOPs), and the capture substrate required to detect, measure, and mitigate hallucinations as they emerge. We argue that AI failure modes are retrospective and emergent, so capture must be comprehensive; anything not captured before it emerges is permanently unavailable to retrospective study. We describe the repository’s reference architecture, its immutable-event data model, a measurement framework, a hallucination risk lifecycle modeled on adverse-event surveillance, and a forensic model for the post-incident investigation of AI harm. We close with a set of dated, falsifiable predictions.

1 Introduction

Organizations are placing generative AI into their business processes faster than they are building the means to manage it. A model now drafts customer replies, writes code, summarizes documents, and helps make decisions, and each use carries both a promise and a hazard. The promise is that AI raises what a workforce can produce. The hazard is that a generative model can be wrong in ways that are fluent and confident, and a wrong answer can move into a business process and cause harm. Governance frameworks and regulation have begun to treat the record of how AI is used as an obligation rather than an option [6, 1]. Most deployments keep no such record.

The specific problem is that an uninstrumented AI deployment cannot be managed. An organization that does not capture how its AI is used cannot measure the return on what it spends, and it cannot prove the value that would justify the investment. It also cannot see the failures. The efficacy of a system prompt or a user prompt is not guaranteed, and a prompt that works today can induce a hallucination under a different input, a longer context, or a changed model. Without a record that hallucination passes undetected and unmitigated, and the harm can propagate and cascade into corporate embarrassment and liability. Value and risk are invisible for the same reason: there is nothing to look at.

In this paper we propose a client architecture that captures AI utilization metadata and session data as the precondition for managing AI in an organization. The architecture is simple at the top and its complexity is a drill-down. The AI systems in use are instrumented, their session content is stored, metadata is derived into a small set of repositories separated by sensitivity, and outcomes and risks are drawn from that metadata and reported to the organization. We argue that capture has to come before the questions it will answer, because

the behaviors worth measuring, from a new failure mode to a new objective, are not known when the data is captured, and anything not captured before it emerges is lost to later study. On this record the harms of AI become foreseeable, detectable, and mitigatable, and the value of AI, including how far it multiplies human capital, becomes calculable. This is a proposal and a theory of operation, not a report of results.

This differs from the tools that instrument AI today. Current operational tooling is largely monitoring: it traces calls and scores them against evaluations chosen in advance, and it serves developers building a single application [5]. Monitoring answers questions its designers already hold. The failures and the value that matter most to an organization are the ones nobody thought to declare, so an enterprise needs an observability substrate rather than a monitoring dashboard, and it needs that substrate to carry governance and continuity as well as measurement. We treat hallucination detection, which a body of work addresses on its own [4], as one risk detector within this larger system rather than as the system itself.

The remainder of the paper is organized as follows. Section 2 sets out the management problem and the use cases that make the record indispensable. Section 3 gives the epistemological argument for capturing before questioning. Section 4 positions the work against observability tooling and governance frameworks. Section 5 describes the architecture and Section 6 its data model. Section 7 develops measurement, portfolio reporting, and the outcome repository. Section 8 treats outcomes and the hallucination taxonomy, and Section 9 the forensic investigation of an AI harm after it occurs. Sections 10 and 11 treat the obligations AI must meet and the trust it must earn. Section 12 describes operation and the planned continuity capability. Section 13 states the paper's predictions, Section 14 its governance, and Section 15 its limits.

2 The Problem: Managing AI Without Visibility

Organizations are placing generative AI into everyday business processes faster than they are building the means to see what it does. A model drafts a contract clause, answers a customer, writes code, or summarizes a case file, and the interaction leaves little behind. The output is used and the conditions that produced it are gone. An organization in this position runs a process it cannot inspect.

The gap is not a single missing dashboard. It is a set of questions an organization cannot answer about its own use of AI. A manager accountable for a deployment would want to establish each of the following and today usually cannot.

Inventory. Which AI tools, models, and instances are in use, by whom, and at what version. Adoption often runs ahead of any record of it, so the inventory itself is unknown.

Utilization and cost. How much the organization uses AI, measured in sessions, tokens, and spend, and how that use is distributed across teams and workflows.

Effectiveness. Whether the AI is doing useful work. Task success, rework, acceptance, and escalation are the measures of performance and effectiveness that would tell a manager whether the deployment earns its cost.

Risk. How often the model produces wrong or fabricated output, whether such failures recur, and whether they cascade into downstream work. Hallucination incidence, and its trajectory over time, are invisible without a record of the interactions that carried them.

Degradation within a session. How context-window utilization and context rot affect output quality as a session lengthens, and when a session has grown long enough that its output can no longer be trusted.

The people and the model. How skillfully the workforce uses the tools, and how the model’s own behavior shifts over time.

The natural response is to add metrics. That response fails for the reason developed in Section 3: the failures that matter are not known in advance, so a fixed set of metrics measures the wrong things. What the organization lacks is not a dashboard but a record. The rest of this paper describes that record and the architecture around it.

2.1 Why the record is indispensable

An organization that puts a generative model into a business process takes on a risk it cannot see without a record. The efficacy of a system prompt or a user prompt is not a given. The same prompt that works today can, under a different input, a longer context, or a changed model, induce a hallucination. That hallucination can pass into the business process undetected, because nothing in an uninstrumented deployment is watching for it, and unmitigated, because nothing catches it. From there the harm can propagate into downstream work and cascade across a chain of decisions, and at the end of that chain are corporate embarrassment and corporate liability.

What makes the absence of the record indefensible is that these events are foreseeable, detectable, and mitigatable. They are foreseeable because prompt efficacy is known not to be guaranteed. They are detectable because the interaction that carried the fault can be examined. They are mitigatable because a detected fault can be corrected before it spreads. An organization that could have foreseen, detected, and mitigated a harm, and did not because it kept no record, owns the consequence. The architecture is what turns a foreseeable, detectable, and mitigatable harm from a realized loss into a managed one. Three cases show the difference.

A prompt that quietly stops working. A system prompt constrains a customer-facing assistant until a customer’s phrasing pushes the model outside it, and the assistant states a policy the company does not hold. Without the record no one knows until a customer acts on it. With the record the interaction is captured, the deviation from the system prompt is detected, and the pattern is caught before it recurs.

A fabrication that enters the record as fact. An analyst asks a model to summarize a source, and the model adds a figure the source does not contain. The figure moves into a report, then into a decision. Without provenance the fabricated figure is indistinguishable from a sourced one. With the record the figure is traceable to a generation with no source, and the correction can reach every place it went.

A cascade across a long session. A session lengthens, the model forgets an earlier constraint, and later outputs violate it while each looks locally reasonable and passes review on its own. Without session capture the drift is invisible. With it, context-window use and the session trajectory are tracked, the drift is detected, and the session is flagged before its later output is trusted.

In each case the difference between a managed event and a corporate loss is the record. The harm was foreseeable, the fault was detectable, and the spread was mitigatable, and only the architecture makes foresight, detection, and mitigation available at the moment they matter.

The same record that makes harm manageable makes value calculable. Without it, an organization can neither prevent the losses of this section nor prove the gains that justify the investment, and it manages its use of AI by anecdote on both counts. Section 7 develops the value side; the point here is that both sides depend on the same record.

3 Epistemological Foundation

The central claim of this paper is about knowledge, not technology. Before an organization can manage its use of generative AI, it has to be able to see that use. We argue that the data required to see it must be captured before anyone knows which parts of it will matter, and that this ordering, capture first and questions later, is forced rather than convenient.

3.1 Two claims about visibility

The first claim is that visibility is a precondition, not a feature. An organization cannot manage, measure, or mitigate what it did not record. Management of an AI deployment is downstream of a record of that deployment. A key performance indicator, a measure of performance or effectiveness, and a hallucination incidence rate are all functions computed over captured data. Where there is no capture, the function has no domain, and the question cannot be asked at all.

The second claim is that the behaviors worth measuring are not known when the data is captured. A generative system produces new failure modes as its models change, as usage shifts, and as context accumulates over a session. A failure mode becomes measurable only after two things happen: enough instances accumulate to form a pattern, and someone builds the lens that renders the pattern legible. Both happen after capture, not before. Capture therefore cannot be conditioned on knowing what to look for. It has to record broadly, in advance of the question.

These two claims meet at an asymmetry. Capture omitted before a phenomenon emerged cannot be recovered after it emerges. There is no retrospective study of data that was never written down. A gap in capture is not a delay that a later effort can close; it is a permanent blind spot in the record. Capturing now is cheap and reconstructing later is impossible, and that asymmetry is the core of the argument.

3.2 Monitoring answers known questions; observability answers unknown ones

The distinction between monitoring and observability makes the argument precise. Monitoring instruments a system against questions its designers already hold: declared metrics, thresholds, and dashboards. It answers known-unknowns. Observability records a system's behavior at enough fidelity and granularity that questions not anticipated at instrumentation time can still be answered later, by interrogating the record. It answers unknown-unknowns [5].

Most current AI operational tooling is monitoring. It traces calls and scores them against evaluations chosen in advance. That is useful and it is not enough. Because AI failure modes emerge, the questions that matter most are the ones nobody thought to declare. A repository built to manage AI has to be an observability substrate rather than a monitoring dashboard. Its job is to keep a future, unspecified question answerable. We call this preserving interrogative optionality, and in Sections 5 and 6 we justify each design choice, immutability, granularity, provenance, and retention, by whether it keeps such a question answerable.

3.3 The limit of retrospection

The power of this approach is bounded by the same fact that motivates it. An organization can study only what capture preserved, and capture is itself a designed, evolving, and imperfect instrument. Three consequences follow, and we state them plainly because they set the boundary of what the repository can promise.

First, a capture schema is a theory of what matters, and it is always partial. A field that is not captured is a question that is foreclosed. Second, capture changes over time. A vendor alters an interface, a field is dropped, a transcript format shifts. Each change draws a line in the record, before which a class of question

is answerable and after which it is not. Third, instrumentation is not free of effect on the system it observes. Recording a session can change how the session is conducted.

These consequences do not weaken the argument. They sharpen its prescription. If capture gaps are permanent and capture is imperfect, then the rational response is to capture earlier, wider, and at higher fidelity than any present question requires, and to treat the capture schema itself as a versioned, first-class artifact whose changes are recorded alongside the data they govern.

3.4 From claim to mandate

The preceding is a claim about how knowledge of an AI system becomes available. It is not yet a policy, and the policy follows as a conclusion. If detectability is retrospective, and if omitted capture is unrecoverable, then an organization that intends to manage, measure, or govern its use of generative AI has no alternative to building the record first. Every downstream capability in this paper, the inventory of Section 5, the measures of Section 7, and the risk lifecycle of Section 8, presupposes it. Emerging governance frameworks and regulation are beginning to treat this record as an obligation rather than a convenience [6, 1]. The remainder of the paper describes the architecture that makes such a record durable, queryable, and safe to keep.

4 Related Work and Positioning

This work sits among three bodies of prior work and departs from each in a specific way.

The first is the operational tooling for large language model applications, often called LLM observability or LLMops. These systems trace calls, record prompts and responses, and score outputs against evaluations chosen in advance, and they serve the developer building and debugging a single application. Their design goal is monitoring in the sense of Section 3: they answer questions their users already hold [5]. This paper takes the same raw material, the captured interaction, and puts it to a different end. The repository is an enterprise-wide observability and governance substrate rather than a per-application debugging tool, and its purpose is to keep answerable the questions an organization has not yet formed.

The second is the body of AI governance and risk frameworks. The NIST AI Risk Management Framework [6], the management-system standard ISO/IEC 42001 [3] and the risk guidance of ISO/IEC 23894 [2], the record-keeping obligations of the EU AI Act [1], and the application risk catalogue of the OWASP Top 10 for large language model applications [7] say what an organization must govern and record. They do not say how to hold the record that makes governance demonstrable. This paper describes that record and argues that it is the substrate those frameworks presume.

The third is the method literature this paper borrows from. Treating hallucination as a detectable, classifiable failure draws on the survey and detection work in natural language generation [4], and the lifecycle of Section 8, from instance to incidence to mitigation efficacy, is adapted from adverse-event surveillance in public health. The contribution here is not a new detector or a new measure but the architecture that lets any of them run over a persistent, governed record.

5 Reference Architecture

The architecture is simple at the top, and its complexity is a drill-down from a few components. At the top it is a flow. The AI systems in use are instrumented. Their session content is stored. Metadata is derived from that content into a small set of repositories separated by sensitivity. Outcomes and risks are drawn from that metadata, and the whole is reported to the organization. Figure 1 shows the flow, and the rest of the paper drills into each part.




Figure 1: The architecture as a simple flow. Instrumented AI systems emit session content to the session content repository, where channels arise and detection, risk, and tagging are enacted. Derived metadata is separated by sensitivity into performance, hallucination-and-risk, and session-continuity repositories. Outcomes and risks are drawn into the outcome and risk repository and presented by outcome and performance reporting. Every record is attributable by organization, product line, and purpose.

5.1 Instrumented systems

The organization runs AI systems, each serving a community of purpose. Every such system is instrumented for utilization: customer journey, throughput, user and session counts, session context-window use, and compaction, among others. This instrumentation is the source of the performance picture. Because each system serves a purpose, its data is attributable to the organization and product line it serves and aggregable by purpose and product. The set of instrumented systems, identified by purpose and product, is the organization’s live inventory of deployed AI.

5.2 Session content repository

The session content repository holds the session data itself: the system and session prompts, the human and model turns, and the session context. Channels arise here, organized by mission and purpose, and a user’s session context supports the several swimlanes that user occupies within a system. Hallucination detection, risk detection, and taxonomy tagging of mission data are enacted on this content, because they operate on what was said, and the content is stored and archived as the system of record.

5.3 Segregated metadata repositories

Metadata is derived from session content into repositories separated by sensitivity, and the separation is deliberate. Performance metadata, the utilization and throughput picture, is one repository. Hallucination and risk metadata is a separate repository, because it is corporate-sensitive and must not be comingled with performance reporting. Session-continuity metadata, the user session data that a later continuity capability would reinject, is segregated again, along boundaries of purpose, product, and organizational sensitivity, because it is mission-centric. Separating by sensitivity, rather than gathering everything into one store, is what lets the sensitive records be governed on their own terms (Section 14).

5.4 Outcome and reporting

The outcome and risk repository draws the positive and negative outcomes associated with AI use, together with the hallucination and risk reporting retrieved from the session-data metadata. It is the store of consequence of Section 7. Above it, outcome and performance reporting retrieves statistics from the metadata

Group	Captured fields
Identity and routing	pseudonymous user id; workflow (lane) id; session id; tenant or organization; timestamp
Model configuration	model id and version; system prompt; decoding parameters (temperature, top- p , max tokens); available tool and function definitions
Interaction content	user prompt (turn); model response; tool calls and results; human interactions (edit, acceptance, rejection, follow-up)
Derived and links	tags; token counts; latency; cost; links to antecedent and consequent events

Table 1: Fields captured per event. Recording the generating conditions, not only the exchanged text, is what makes an event explicable after the fact.

repositories and presents organizational and corporate performance. This is the dashboard and the overall performance reporting system, reporting against the layered objectives of Section 7 and attributing performance to organizations and product lines.

Each part is a drill-down. The lower parts preserve and organize the record, the upper parts interrogate and report it, and Section 3 argued that the record has to come first and be kept whole.

6 Data Model

The data model follows from the epistemology of Section 3. If an instance’s meaning is knowable only through its cause and its trajectory, then the unit of capture cannot be a flat log line. It has to be a first-class, immutable event that carries enough context to be explained after the fact and linked to what came before and after it.

6.1 The captured event


Each interaction between a user and a generative model is stored as one immutable event. An event records the exchanged text and the conditions that produced it, because a response is interpretable only against the instruction frame and the sampling regime that generated it. A hallucination produced at a high sampling temperature under a permissive system prompt is a different event, with a different cause, from the same surface text produced at temperature zero under a constraining one. Recording the output without its generating conditions discards the causal information the repository exists to preserve.

Table 1 lists the captured fields. The user identifier is a stable pseudonymous handle, not personal data. It supports per-user measures of utilization and acumen (Section 7) and workplace monitoring of system use, while the repository avoids storing personal identifiers. Section 14 describes the governance that keeps identity trackable to a role and a workflow without retaining personally identifiable information. Every field is tagged and searchable, so events can be retrieved by user, workflow, model, system prompt, sampling regime, or content.

6.2 Immutability, provenance, and links

Events are immutable and versioned. A later re-classification of an event adds a labeled interpretation without altering the original record, so a lens built after capture can be applied without destroying the data it reads. Each event carries directed links to its antecedents and consequents, so a chain of exchanges, and a cascade of failures, can be traversed rather than inferred. Section 8 uses these links to trace recurrence and causality.

Two forms of each record



Channels, linked




Figure 2: The data model. Each record is kept in two forms: a raw, append-only form that is never rewritten, and a structured, queryable form derived from it, so a later re-classification adds a label without altering the original. Records are organized into typed channels (event, outcome, risk, hallucination) and linked, both along the event channel as antecedents and consequents and across channels, so a derived record can be read together with the event it arose from. The fields captured per event are listed in Table 1.

The repository keeps each record in two forms. A raw, append-only form preserves each interaction exactly as captured and is never rewritten. A structured form makes the fields of Table 1 queryable at scale and carries the links between records, so a chain of exchanges, and a cascade of failures, can be followed by query. Linkage is a property of the records, not a separate store. Figure 2 shows the two forms and how records link across channels.

6.3 Channels

The session content repository is not a single stream. It is a set of channels, each a typed, linked stream of one kind of record, organized by mission and purpose. Session content, detections, and outcomes are distinct channels that link to one another, so a risk can be read together with the event that raised it and the outcome it threatens. Which repository a channel is stored in is a matter of sensitivity, not model. Session content is archived in the session content repository, while derived metadata is segregated into the metadata repositories of Section 5 by purpose, product, and sensitivity. What the model fixes is that every channel is typed and linked, so segregating the stores never severs the links between records.

7 Measurement Framework

Measures of performance, measures of effectiveness, and key performance indicators are not properties of the AI architecture. They are the strategic and operational objectives of the organization, defined by the organization and by the corporation above it, in layers. A corporation sets strategic objectives, a division sets operational objectives that serve them, and a team or workflow sets local objectives in turn. Each layer names its own measures. The architecture does not define these measures and should not. Its role is to collect, detect, catalog, and store the outcomes those measures are computed from, so that each layer can report against its own objectives. Figure 3 summarizes the framework.




Figure 3: The measurement framework. Objectives are exogenous and layered: each of the workflow, organizational, and corporate layers defines its own measures. The architecture does not set the measures; it computes them over the session and outcome record and rolls outcomes up the hierarchy. Measures of performance sit close to the interaction, measures of effectiveness against the objective a workflow serves, and key performance indicators at the business-facing top, alongside the workforce and model acumen measures, the trust-calibration gap, and customer experience.

7.1 Objectives are defined by the organization, in layers

The organization owns the objective hierarchy, and the architecture holds a model of it: a catalog of the objectives each layer has declared. A captured interaction is associated with the objectives it bears on, and outcomes roll up from the workflow where they occur to the strategic objective they serve. Because the objectives are exogenous, the architecture adapts to whatever an organization declares rather than imposing a fixed set of metrics. This is the principle of Section 3 applied to measurement: the record has to support measurement against objectives that were not anticipated when it was captured, because an organization revises its objectives over time.

7.2 The outcome repository

An outcome is a realized value of a measure, attributable to AI use, at a point in time. The architecture treats outcomes as first-class records, stored in an outcome repository alongside the event record of Section 6. Outcomes are collected where they are declared, detected where they are implicit, since a resolved case, an accepted draft, or a completed transaction is an outcome the record can recognize without being told, cataloged against the objective hierarchy, and stored so that reporting draws on a persistent, queryable history rather than a snapshot. Because each outcome links to the interactions that produced it, an outcome can be traced to the AI use it is credited to or charged against.

The repository stores outcomes of both signs together with the conditions around them. It holds positive outcomes and negative outcomes, and three classes a measurement-only store would omit: risks, which are anticipated negative outcomes not yet realized; violations, which are breaches of the obligations of Section 10; and insider-threat indicators, which are signs that AI is being used to act against the organization.

Each is a consequence, realized or potential, that leadership has to see and report, so each takes the outcome record form and links back to the interactions it arose from.

Two dimensions organize the store. The first is organizational structure: an outcome is filed against the unit it belongs to, so that the overall organization and its internal units stay distinguishable and a unit's outcomes can be read apart from the whole. The second is the objective it bears on, from the layered hierarchy above, so that an outcome rolls up from the workflow where it occurred to the strategic objective it serves. Together the two dimensions let a report be drawn for the whole organization or for any unit, against that level's own objectives.

The repository also records the context an outcome arose in. Outcomes from live operation and outcomes from exercises, drills, and simulations are kept apart, so that an organization can rehearse against the record and assess its readiness without contaminating the measures of live performance. Insider-threat and violation records carry the same separation, since a rehearsed exercise and a real breach must never be counted as the same event. The governance of who may read these records, and the privacy limits on monitoring for them, is the subject of Section 14.

7.3 Performance, effectiveness, and indicators

Within this frame the three families sit at different layers. Measures of performance describe whether the system did the thing: latency, availability, token use, and cost, computed close to the interaction. Measures of effectiveness describe whether the objective was achieved: task success, rework, acceptance, escalation, and time to deliver, computed against the objective a workflow serves. Key performance indicators are the business-facing roll-ups each layer reports. To these the architecture adds measures the record makes possible and a fixed dashboard would miss: workforce acumen, computed from how skillfully users direct and check the AI; model acumen and drift, computed from the model's reliability over time; the trust-calibration gap of Section 11; customer experience for customer-facing work, in satisfaction, resolution, and escalation; and token cost as the denominator of return. Measurement is neutral to the sign of what it finds. The same record yields the positive and negative outcomes of Section 8.

7.4 Portfolio reporting

The AI architecture is a funded investment, and like any portfolio investment it has to justify its cost. The measurement layer supports reporting at the level a portfolio owner needs: what the organization spends on AI, set against the outcomes attributable to it, reported against the objectives each layer defined. Cost is known from the record of Section 6. Outcome value is known from the outcome repository. Return is the relation between them, reported at each layer of the objective hierarchy. Without the record, this report is an anecdote. With it, it is an account.

Because the record makes cost and outcome both measurable, the return on the AI investment is calculable rather than asserted. The investment is measurable, the deployment is instrumented, and the achievement of the workforce and the evolution of AI's benefit over time are trackable against the objectives each layer set. The question an executive most wants answered, how far AI multiplies human capital, becomes a computed quantity: the output a workforce produces with AI, set against what it produced without, tracked as it changes. Without the record none of these is calculable, and the investment is defended by anecdote.

7.5 Competency and mission, in both directions

The measurement establishes how the AI architecture affects the organization's competency and its mission deliverables, and the effect runs in both directions. AI can raise organizational competency by completing work faster or better, and it can lower it by displacing skills a workforce would otherwise maintain. A

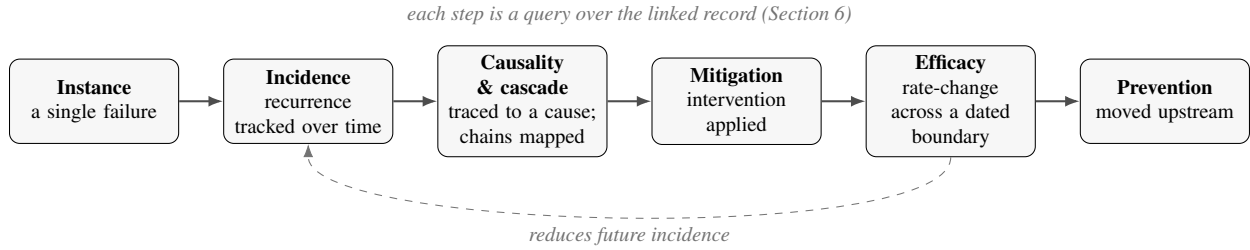


Figure 4: The hallucination lifecycle, adapted from adverse-event surveillance. A single instance recurs into a tracked incidence, is traced to a cause and its cascade, and draws a mitigation whose efficacy is measured as a change in incidence across the dated boundary at which it was applied; an effective mitigation moved upstream becomes prevention and reduces future incidence. Each step is a query over the linked record, not an inference from a point-in-time log.

workforce that defers to the model on a class of task can lose the ability to perform or check that task, which is a degradation of competency even when short-term output looks improved. The measurement layer has to show both, because a portfolio owner deciding whether the investment pays needs the losses as well as the gains. This connects to the trust of Section 11: over-reliance is the mechanism by which competency degrades, and it shows up in the record as reliance that outruns verification.

8 Outcomes: Benefit, Harm, and Liability

An AI deployment produces outcomes of both signs, and the record captures them with the same machinery. On the positive side are mission benefit and performance improvement: work completed faster, at lower cost, or at higher quality than before. On the negative side are error, rework, harm, and liability. Measurement (Section 7) is neutral to the sign of an outcome. What makes an outcome legible is that the interaction which produced it, and the downstream use of its output, are on the record.

8.1 Positive outcomes

Performance improvement and mission benefit are measured as changes in the measures of effectiveness of Section 7, attributed to AI use through the record: which workflows use AI, how their task success, cost, and time change, and whether the change holds. Without the record, benefit is asserted anecdotally. With it, benefit is measured and attributed to the work that produced it.

8.2 Harm and the adverse-event lifecycle

Negative outcomes follow a lifecycle borrowed from adverse-event surveillance. A single failure is an instance. Its importance is not evident at the instant it occurs, because impact is a function of trajectory. Repeated instances form an incidence that can be tracked. Some instances trace to a cause, and some cascade into further failures downstream. A mitigation applied after detection has an efficacy that is itself measurable, as a change in incidence across a dated boundary, and a mitigation moved upstream becomes prevention. Each step is a query over the linked record of Section 6, not an inference from a point-in-time log. Figure 4 shows the lifecycle.

Category	How it appears in use	Repository signal
Fabrication	Asserts a fact, source, feature, or justification that is not true or not grounded	Output checked against provenance; assertions with no grounded source
Overconfidence and misrepresentation	Asserts without verification, or reports a task as done or correct when it is partial or untested	Confidence and completion claims vs verification steps and downstream outcome
Staleness	Treats outdated information as current, or omits a freshness check	Cited sources and state vs their capture time and current values
Context loss	As a session lengthens or crosses a boundary, drops, reshapes, or forgets earlier content	Session trajectory and context-window use; later output vs earlier record
Attribution error	Misassigns a source, an authority, or a capability	Claimed attributions vs provenance links
Scope and instruction divergence	Does more than, or other than, what was asked	Output scope vs the recorded prompt and workflow boundary
Agentic and tool-use failure	Acts without verifying preconditions, misuses or ignores a tool, references a resource that does not exist, or takes an unsafe action	Tool calls, preconditions, and referenced resources vs the recorded environment

Table 2: A working taxonomy of hallucination the repository is designed to characterize and detect. Categories are not mutually exclusive; a single event can carry more than one. Applicable categories depend on how the organization uses AI.

8.3 A working taxonomy of hallucination

Harm has structure. The failures grouped under the word hallucination are not one thing, and detecting them requires knowing which form is in play. The repository classifies hallucination against a working taxonomy, summarized in Table 2. The taxonomy matters to the architecture in a specific way. Each category is detectable only from particular captured fields, so the taxonomy is the reason the data model of Section 6 records the generating conditions, the provenance, and the session trajectory, and not the output alone. The categories synthesize forms recognized in the hallucination-detection literature [4] with operational forms specific to deployed and agentic systems.

Two failure modes travel alongside hallucination without being instances of it, and the record surfaces both. The first is refusal and over-restriction, where the model declines a legitimate request or narrows a permitted one, so the harm is a withheld benefit rather than a wrong output. The second is its inverse, compliance with a request the model should have declined, where the harm is an action taken that a boundary should have stopped. Both are legible only against the recorded request and the workflow’s permitted scope, and both are measured the same way the outcomes of Section 7 are: over-refusal as a loss of effectiveness, under-refusal as a violation of the obligations of Section 10.

Fabrication is treated at greater length next, because it is the form most damaging to an organization’s fact base.

8.4 Fabrication masquerading as fact

The sharpest harm is specific to generative systems. A hallucination is not only a wrong answer. It is fiction presented in the form of established fact, fluent and confident, and once it enters an organization’s records it can corrupt the fact base that later work depends on. The architectural defense is provenance. Because every event carries the conditions and source of its content (Section 6), a generated assertion stays distinguishable from a curated fact. The repository does not prevent fabrication at the moment of generation. It prevents

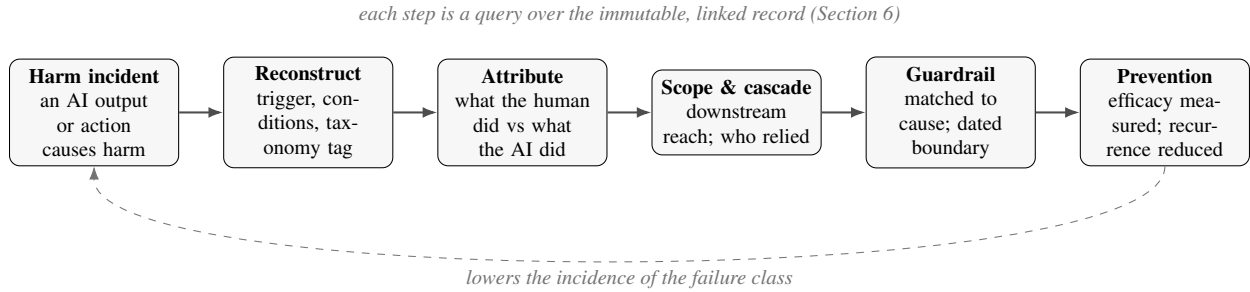


Figure 5: The harm-forensics loop. A harm incident opens an investigation that reconstructs the trigger and conditions, attributes the event between the human actor and the AI actor, and traces its downstream scope and cascade. The investigation prescribes a guardrail matched to the cause it found, applied at a dated boundary, whose efficacy is measured as a change in incidence, so an effective guardrail moved upstream becomes prevention and lowers the recurrence of the failure class. Each step is a query over the immutable, linked record of Section 6.

fabrication from silently becoming fact, by preserving the boundary between what a model asserted and what was established as true.

8.5 Liability

Negative outcomes carry liability, and liability attaches to specific acts. Establishing what the model was asked, under what configuration, what it produced, who relied on it, and what followed requires exactly the record this paper describes. The obligations that an outcome may violate, and against which liability is assessed, are the subject of Section 10.

9 Harm Forensics and the Post-Incident Investigation

When an AI-induced harm occurs, the harm is not the end of the event. It opens an investigation the organization is obliged to conduct, and the questions that investigation asks are forensic: what was produced, what triggered it, what caused it, how far it spread, who and what is accountable, and how a recurrence is prevented. None of these can be answered after the fact for an interaction that left no record. Among its other roles, the architecture of this paper is the forensic substrate for AI harm. The immutable, linked, provenance-bearing record of Section 6 is what makes a post-incident investigation possible at all, and the adverse-event lifecycle of Section 8 is what turns the investigation’s findings into a defense against the next incident. Figure 5 shows the loop this section describes.

9.1 The forensic questions and where the record answers them

A post-incident investigation has to establish a specific set of facts, and each is a property of the captured record rather than an inference drawn after the event. Table 3 lists the questions a forensic account of an AI harm has to answer and the evidence in the record that answers each. The point is not that these facts are easy to recover, but that they are recoverable at all only where the interaction was captured before the harm was known.

Forensic question	Evidence in the record
What was produced, and was it wrong	The model response, tagged against the hallucination taxonomy (Section 8); a fabricated assertion distinguished from a grounded one by provenance
Under what conditions it was produced (the trigger)	The generating conditions of the event: system prompt, decoding parameters, input, and model id and version (Table 1)
Was a hallucination or fabrication involved, and of which kind	The taxonomy tag on the event and the captured fields that support each category
What the human actor did	The human turns and the recorded human interactions: edit, acceptance, rejection, and follow-up
What the AI actor did	The model turns, the tool calls, and the tool results
How far it spread (scope and scale)	The consequent links from the event, traversed to the downstream work the output entered
Whether a control existed and held	The presence or absence of a review, rejection, or verification step in the recorded interaction

Table 3: The questions a forensic investigation of an AI harm must answer, and the evidence in the record that answers each. Each answer is a query over the captured event and its links, not a reconstruction attempted after the interaction is gone.

9.2 Separating the human actor from the AI actor

A forensic account has to attribute the harm, and attribution begins by separating what the human actor did from what the AI actor did. The record supports this because it captures both sides as distinct, linked events: the human turns and the human interactions of Table 1 on one side, and the model turns, tool calls, and tool results on the other. From this an investigation can establish whether a harm originated in the model’s output, in a person’s use of that output, or in the interaction between them, and whether a human review gate existed at the point of harm and whether it held.

This separation is what distinguishes two adversarial cases an investigation must not conflate. In one, the AI actor is the source: the model fabricates, misrepresents, or takes an unsafe autonomous action, and the harm is a model failure of the kind Section 8 classifies. This is the rogue or agentic case, and its signature in the record is an action the model took without a verified precondition or outside its authorized scope. In the other, the human actor is the source: a person directs the AI to act against the organization, and the harm is insider misuse for which the AI is the instrument. Its signature is not a model error at all but a pattern of instruction and use, such as extraction at volume, probing of scope or permission, or direction toward a prohibited end. The same record grounds both findings, and only the record lets an investigation tell them apart, because the intent of a human actor is legible only through the sequence of instructions and uses the record preserves.

9.3 Scope, scale, and the cascade

The impact of an AI harm is a function of how far its output travelled, and that reach is rarely evident when the harm occurs. A fabricated figure enters a report, the report enters a decision, and the decision conditions work downstream of it. Without the links of Section 6 an investigation sees only the first event and underestimates the harm. With them it traverses the consequent links to establish scope, which work was affected, and scale, how much, so a correction can reach every place the output went. This is the forensic form of the claim registered as Prediction 4: absent the record, the impact of an incident is systematically underestimated, because the cascade linking it to its consequences is not observable at the moment it occurs.

Guardrail class	What it changes	The cause it addresses
System-prompt engineering	Constrains the instruction frame the model operates under	A harm traced to a permissive or ambiguous system prompt
Input and output validation	Checks inputs and screens outputs before they are used	A harm traced to an unchecked input or an unscreened assertion
Human-in-the-loop review gate	Requires human review at a defined point before an output is used or an action is taken	A harm a review gate would have caught, or one where a gate existed but was not enforced
Tool and permission constraint	Limits the actions an agentic system may take without authorization	A harm traced to an unsafe or unauthorized autonomous action
Workforce hallucination-awareness training	Raises the workforce’s ability to recognize and check AI failure	A harm traced to over-trust or an unrecognized fabrication (Section 11)
Model version and configuration control	Governs which model versions and settings are deployed	A harm traced to a model change or a configuration drift

Table 4: Classes of guardrail a post-incident investigation can prescribe, each matched to the class of cause it addresses. Each guardrail is itself recorded, applied at a dated boundary, and evaluated by the same measurement the harm was detected with, so a response whose efficacy the record cannot demonstrate is distinguishable from no response.

9.4 From post-mortem to guardrail

A forensic investigation is not complete when it has explained a harm. Its purpose is to prevent the next one, and its output is a guardrail. The lifecycle of Section 8 supplies the mechanism: a mitigation is applied at a dated boundary, its efficacy is measured as a change in incidence across that boundary, and a mitigation that proves effective and is moved upstream becomes prevention. What the post-mortem adds is the choice of guardrail, informed by the cause it found. Table 4 lists the classes of guardrail a post-mortem can prescribe, each paired with the cause it addresses.

Because each guardrail is recorded and dated, the organization can show not only that it responded to a harm but that its response worked, and can tell an effective guardrail from a nominal one. A guardrail whose efficacy the record cannot demonstrate is, for the purpose of accountability, indistinguishable from no guardrail at all.

9.5 Forensics and liability

The forensic record is also the basis on which liability is assessed. Liability attaches to specific acts (Section 8), and establishing what the model was asked, under what configuration, what it produced, who relied on it, and what followed is the reconstruction this section describes. An organization that can produce this account can bound its liability to what occurred and demonstrate the diligence of its response. An organization that cannot is exposed to the worst construction of an event for which it holds no record.

10 Obligations: What AI Must Be Evaluated Against

An AI output or action is not evaluated against correctness alone. It is evaluated against the obligations that bind the person or organization that used it. An answer can be factually right and still violate a policy, a law, a professional standard, or a required procedure. To manage AI, an organization has to evaluate its use against each class of obligation that applies, and that evaluation is possible only over a record that holds the output together with its conditions and its downstream use.

Class	What it obliges	Example failure
Organizational	Internal policy, brand and communication standards, role and workflow scope	Acting outside an authorized workflow
Legal and regulatory	Statute, regulation, contract and licensing terms, intellectual property, data protection	Reproducing licensed text; mishandling regulated data
Ethical	Fairness, non-deception, non-maleficence, transparency, avoidance of bias	Biased screening; undisclosed automation
Procedural	Required process, review gates, verification and citation, documentation	Skipping a mandated human review or source check
Professional and domain	The standard of care and competence of the profession or domain	Advice below a professional standard
Factual and episodic	Grounding in established fact, provenance, non-fabrication	Fabrication presented as fact (Section 8)
Safety and operational	Avoidance of harmful action, security, reliability	An unsafe automated action

Table 5: A sample set of obligation classes an AI use can be evaluated against. The set is not exhaustive; applicable classes depend on the domain.

Table 5 gives a sample set of obligation classes. The set is not exhaustive, and which classes apply depends on the domain, but the classes are general enough to frame the evaluation for most deployments.

The repository makes each of these evaluations possible for the same reason. Conformance and violation are both properties of a recorded interaction and its consequences, and neither can be assessed for an interaction that was never captured. Evaluation against obligations is a use of the record, and a violation is one of the negative outcomes of Section 8.

11 Trust as an Earned Property

Trust is central to AI deployment, and it is the property most likely to be misplaced. Users tend to treat an AI as an expert. Unless they are trained otherwise, they attribute expertise and authority to it that it has not earned, and they extend trust its record does not justify. Fluent, confident output invites this. The result is unwarranted trust, and it is dangerous precisely because the system’s failures, as Section 8 described, arrive in the form of confident fact.

Trust should instead be calibrated. Reliance on the AI should track its demonstrated reliability. Trust that exceeds reliability is over-trust, and it turns each fabrication into an accepted fact. Trust that falls short of reliability is under-trust, and it forfeits benefit the AI could provide. The goal is not maximal trust or minimal trust but reliance that matches evidence.

Calibration requires two things the repository provides. Trust has to be earned, which means measured: the reliability of the AI on a task, a workflow, or a class of question is a statistic computed over the record, not an impression. And trust has to be curated, which means surfaced: the provenance of an output, the model’s track record on similar work, and the confidence attached to an assertion have to be presented to the user, so that reliance tracks evidence rather than fluency. A user who can see that an assertion is a generated claim without a source, on a task where the model has a poor record, is a user who can withhold trust that has not been earned.

Managed this way, trust becomes a property the deployment measures and cultivates rather than a disposition users bring by default. This connects the workforce acumen of Section 7, since a habitually over-trusting user is a measurable risk, to the obligations of Section 10, since misplaced trust is how an obligation violation reaches a consequence unchecked.




Figure 6: Client deployment. Instrumentation at the AI-use surfaces, an API gateway or proxy, hooks in agent frameworks, and connectors to chat and development interfaces, captures each interaction and collects it to the repository, which is then aggregated and served as reporting. A health-and-availability monitor watches capture, storage, and reporting, and feeds both operational monitoring and value modeling, because an observability system has to be observable.

12 Theory of Operation

The architecture is deployed as a client-side system that captures at the points of AI use, stores to the repository, and serves reporting back to the organization. This section describes the functional components it runs and the topology it is deployed in, and then the continuity capability planned for a later phase. Figure 6 shows the deployment.

12.1 Functional components

Three functional components run across the parts of Section 5. The first collects and stores: it ingests each interaction and outcome and persists it to the channels of Section 6. The second aggregates and reports: it rolls records up the objective hierarchy and the organizational structure and serves the reports of Section 7. The third monitors health and availability: it watches whether capture, storage, and reporting are themselves working, because an observability system has to be observable. Its output serves two ends. It drives operational monitoring, so that an outage in capture is seen rather than mistaken for an absence of activity, and it feeds value modeling, since availability is part of the value an organization receives and downtime is value it does not.

12.2 Deployment topology: convergence and isolation

The channels of Section 6 can be converged into one repository with logical separation, or split into isolated systems, and the choice is a matter of assurance. Where separation is a convenience, convergence with the context and structure dimensions of Section 7 is enough. Where separation has to be guaranteed, the systems are isolated rather than converged. Live operation and exercises can run on separate systems, so that a drill cannot spill into live measures and a live incident cannot enter a rehearsal. A sensitive organizational product line can run on its own isolated system, so that its record cannot contaminate, or be contaminated by, the rest. Isolation trades the convenience of a single store for a guarantee that certain records never meet,

and an organization chooses, per population, which guarantee it needs.

12.3 Future leg: controlled continuity and workflow isolation

A later phase of the system adds two capabilities that share one principle: context should move between sessions only under explicit control.

The first is continuity reinjection. Because a session’s record is preserved in full, a later session can be resumed by re-injecting a curated summary of the prior one, so that work continues across the session boundary rather than restarting at it. The reinjected context is drawn from the record, is itself recorded, and is scoped to the workflow it belongs to.

The second is workflow isolation. A single user runs several distinct workflows, and these must not intermingle. The system gives each user a set of isolated lanes, and a session is bound to one lane. Continuity reinjection draws only from the lane’s own history, and one lane’s context cannot leak into another. Isolation is not only an organizational convenience. Uncontrolled mixing of unrelated context is itself a source of error in generative systems, because a model conditioned on material from an unrelated task will draw on it. Binding a session to a lane makes continuity deliberate and bounds the context a response can be conditioned on.

Together these turn session continuity from an accident of whatever remains in a context window into a governed operation. Continuity is restored on purpose, from the record, within a boundary.

13 Registered Predictions

This is a proposal, and its claims will be tested by deployment rather than by results reported here. To make that test meaningful, we state the paper’s central predictions now, as specific and falsifiable claims, so that later observation in deployed systems can confirm or refute them against the dated record. Each prediction follows from the theory of operation and is checkable from the kind of data the architecture captures.

1. Hallucination recurrence for a given type clusters in time around context-window saturation and compaction events, rather than arriving at random.
2. In a corpus captured broadly, the rate at which newly detectable failure classes appear rises with the age and size of the corpus, because detection is retrospective and emergent (Section 3).
3. The efficacy of a mitigation is measurable as a change in incidence across the dated boundary at which it was applied, and the effect either persists or decays measurably over the sessions that follow.
4. Without a repository, the impact of an incident is systematically underestimated, because the cascade linking it to its consequences is not observable at the moment the incident occurs.
5. Workforce acumen, measured from correction cadence and instruction specificity, correlates with realized measures of effectiveness independently of the capability of the model in use.
6. A change to the capture schema creates a permanent retrospective blind spot, observable afterward as a discontinuity in a metric that was trackable before the change.
7. When an AI harm is investigated over the record, the reconstructed scope of its cascade exceeds the scope visible at the moment of the incident, and the gap widens with the number of downstream steps the output passed through.

8. A guardrail prescribed by a post-incident investigation and applied at a dated boundary produces a measurable change in the incidence of the failure class it targets, and the change is larger for a guardrail placed upstream of the cause than for one placed at the point of harm.
9. Over-refusal and under-refusal move in opposite directions under a given system prompt, so that a change tightening the instruction frame lowers one while raising the other, observable in the record as a trade-off rather than a uniform improvement.

We record these as of the date of this paper, so that their confirmation, if it comes, is a matter of record and not of hindsight.

14 Governance, Privacy, and Compliance

The record described in this paper is powerful because it is complete, and it is sensitive for the same reason. Governance is what makes keeping it legitimate, and in this architecture governance is not bolted on after the fact. It is expressed through choices already made: the segregation of stores by sensitivity, the use of a pseudonymous identity, and the immutability of the record. This section states how those choices, with a few added controls, keep the record lawful and safe to hold.

14.1 Segregation as the primary control

The separation of metadata by sensitivity (Section 5) is the first governance control, not a convenience. Performance metadata, hallucination and risk metadata, and session-continuity metadata are held apart so that each can be governed on its own terms: its own access list, its own retention, and its own reason for existing. Segregation limits the blast radius of any single access and enforces need-to-know, so that a person who reports on performance does not thereby gain the organization's risk and insider-threat records.

14.2 Identity: pseudonymous by default

The user identifier is a stable pseudonymous handle, not a name. It supports per-user acumen (Section 7) and the monitoring of system use without storing personal identifiers. Re-identifying a handle to a named person is a governed exception, permitted for a stated purpose such as a security case, performed under authority, and itself recorded. This is pseudonymization in the sense data-protection regimes intend: identity is recoverable under control, not absent and not exposed. The default is that the organization measures roles and workflows, not people.

14.3 Monitoring, insider threat, and proportionality

Because the record captures how AI is used, it is also security telemetry, and the insider-threat and violation detection of Section 7 reads it. This capability is legitimate and it is bounded. It is bound by purpose, so that records gathered to measure performance are not repurposed for surveillance without cause. It is bound by access, so that those who read security records are not those who run performance reporting. And it is bound by proportionality, because monitoring is not free of effect on the monitored. Surveillance changes behavior, the observer effect of Section 3 in human form, and an organization that watches too closely will change the work it meant only to measure. The controls that make monitoring lawful are also what keep it from hollowing out the trust it depends on.

14.4 Retention, minimization, and audit

Two requirements pull against each other, and segregation is what reconciles them. Retrospective study (Section 3) wants the record kept long and whole. Data protection wants personal data minimized and kept no longer than its purpose requires. Because the stores are segregated they can be governed differently: the analytic record can be retained for study with personal fields minimized or pseudonymized, while the sensitive stores carry shorter retention and tighter access. Access is granted by role and by store, and because the record is immutable, every access to a sensitive store is itself an entry in the record, so the system that audits the organization's use of AI audits its own use in the same way.

14.5 Standards conformance

The architecture is built to sit inside the frameworks an enterprise is measured against, because conformance is what makes it deployable. It provides the measure and manage substrate that the NIST AI Risk Management Framework calls for [6], and the evidence base an AI management system under ISO/IEC 42001 [3] has to maintain, with risk handled along the lines of ISO/IEC 23894 [2]. It answers the record-keeping and logging obligations of the EU AI Act [1]. Its risk taxonomy (Section 8) overlaps the application risks catalogued in the OWASP Top 10 for large language model applications [7]. The claim is not that the repository satisfies these frameworks on its own, but that it supplies the record without which conformance to them cannot be shown.

Governance and the record are in tension only when governance is an afterthought. Here segregation, pseudonymity, and audit are properties of the architecture, so holding the record and governing it are the same act.

15 Limitations and Threats to Validity

The approach has limits, and they follow from the same facts that motivate it.

The power of the record is bounded by the completeness of capture, and capture is a designed, imperfect instrument. A field not captured is a question foreclosed, and a schema is always a partial theory of what will matter (Section 3). The architecture reduces this by capturing broadly and versioning the schema, but it cannot recover what was never recorded.

Capture also changes over time. A vendor alters an interface, a format shifts, or a field is dropped, and each change draws a line in the record before which a class of question is answerable and after which it is not. The architecture records schema changes as first-class events so that these discontinuities are visible, but visibility is not repair.

Instrumentation is not free of effect on what it observes. Recording a session, and still more monitoring it for risk, can change how the work is done (Section 14). The controls of Section 14 bound this effect but do not remove it.

Finally, this paper is a proposal and a theory of operation. Its claims are argued from the design and stated as predictions (Section 13); they are not demonstrated by results here. The evidence for them will come from deployment, and until then the case rests on the coherence of the argument and the falsifiability of the predictions rather than on measured outcomes.

16 Conclusion

An organization cannot manage what it cannot see, and it cannot see its use of AI without a record of that use. We have proposed a client architecture that captures AI session data and derived metadata as that

record, organized as a simple flow: instrumented systems, a session content repository, metadata repositories segregated by sensitivity, an outcome and risk repository, and reporting. We have argued that capture must precede the questions it will answer, that on such a record the harms of AI become foreseeable, detectable, and mitigatable, and that its value, including how far it multiplies human capital, becomes calculable. The same record serves both ends, so an organization that keeps it can both defend against the liability of AI and prove its return, while an organization that does not can do neither. The predictions of Section 13 state what we expect a deployed system to show. We offer them now, on the record, so that what reads today as a proposal can later be read against what was observed.

References

- [1] European Parliament and Council of the European Union. Regulation (EU) 2024/1689 of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act), 2024. Official Journal of the European Union, L series, 12 July 2024.
- [2] ISO/IEC. ISO/IEC 23894:2023 – Information technology – Artificial intelligence – Guidance on risk management, 2023.
- [3] ISO/IEC. ISO/IEC 42001:2023 – Information technology – Artificial intelligence – Management system, 2023.
- [4] Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Andrea Madotto, and Pascale Fung. Survey of Hallucination in Natural Language Generation. *ACM Computing Surveys*, 55(12):1–38, 2023.
- [5] Charity Majors, Liz Fong-Jones, and George Miranda. *Observability Engineering: Achieving Production Excellence*. O’Reilly Media, 2022.
- [6] National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Technical Report NIST AI 100-1, U.S. Department of Commerce, January 2023.
- [7] OWASP Foundation. OWASP Top 10 for Large Language Model Applications (2025), 2025. OWASP GenAI Security Project; released November 2024.